

Session 11: Missing Data

1. What is a NaN Value?

```
[1]: from numpy import nan      #stands for Not A Number
      nan==3
```

False

```
[2]: nan!=3
```

True

```
[3]: nan is nan
```

True

```
[4]: nan==nan
```

False

```
[5]: nan+3
```

nan

```
[6]: 2*nan
```

nan

```
[7]: sum([nan,3])
```

nan

```
[8]: import pandas as pd
      pd.Series([nan,3,2]).sum()
```

5.0

```
[9]: pd.Series([nan,3,2]).sum(skipna=False)
```

nan

```
[10]: import pandas as pd
       pd.isnull(nan)
       # Opposite is pd.notnull
```

True

```
[11]: pd.isnull('nan')
```

False

Q1: Predict the output of the each of the following lines without typing them, and then verify by typing.

```

nan==True
nan==False
nan**2
nan!=nan
pd.isnull('null')
pd.isnull(nan==True)

```

Answer: The first two statements are False because nan is not equal to anything.

The third statement evaluates to nan because any numerical operation with nans evaluates to nan.

The fourth statement (perhaps surprisingly) is True, because nan==nan is False.

The last two statements are all False. This is because 'null' is a string and 'nan==True' evaluates to False, both of which are not nan.

Q2: Explain the output of the following code using your knowledge of NaNs.

```

[12]: df=pd.DataFrame([[nan,10],[50,nan],[nan,20],[20,10]],columns=['Capacity','Demand'])
      df

```

	Capacity	Demand
0	NaN	10.0
1	50.0	NaN
2	NaN	20.0
3	20.0	10.0

```

[13]: (df['Capacity']-df['Demand']).sum()

```

10.0

Explanation: This takes the difference of the two columns first, and only the last item is not null. Hence, the total is 10. (See below)

```

[14]: df['Capacity']-df['Demand']

```

0	NaN
1	NaN
2	NaN
3	10.0

dtype: float64

```

[15]: df['Capacity'].sum()-df['Demand'].sum()

```

30.0

Explanation: This takes sums the none-null entries in each column first. The result is 70-40, which is 30.

```

[16]: df.sum(axis=0)

```

Capacity	70.0
Demand	40.0

dtype: float64

```

[17]: df.sum(axis=1)

```

```

0    10.0
1    50.0
2    20.0
3    30.0
dtype: float64

```

```
[18]: df.sum(axis=1, skipna=False)
```

```

0    NaN
1    NaN
2    NaN
3    30.0
dtype: float64

```

Explanation: The first command sums the columns and the second sums the rows. The default behavior is to skip NA values. The third command does not skip NA values, so the sum for the first three rows is NaN.

2. Handling Missing Values

Loading the Ebola Dataset

```

[19]: import pandas as pd
      base='https://raw.githubusercontent.com/chendaniely/pandas_for_everyone/master/data/'
      filename='country_timeseries.csv'
      ebola=pd.read_csv(base+filename)
      ebola['Date']=pd.to_datetime(ebola['Date'])
      ebola.set_index('Date', inplace=True)
      ebola.iloc[:5,:6]

```

	Day	Cases_Guinea	Cases_Liberia	Cases_SierraLeone	\
Date					
2015-01-05	289	2776.0	NaN	10030.0	
2015-01-04	288	2775.0	NaN	9780.0	
2015-01-03	287	2769.0	8166.0	9722.0	
2015-01-02	286	NaN	8157.0	NaN	
2014-12-31	284	2730.0	8115.0	9633.0	

	Cases_Nigeria	Cases_Senegal
Date		
2015-01-05	NaN	NaN
2015-01-04	NaN	NaN
2015-01-03	NaN	NaN
2015-01-02	NaN	NaN
2014-12-31	NaN	NaN

Counting Missing Entries

```
[20]: ebola.iloc[:5,:6].isnull()
```

	Day	Cases_Guinea	Cases_Liberia	Cases_SierraLeone	\
Date					
2015-01-05	False	False	True	False	

2015-01-04	False	False	True	False
2015-01-03	False	False	False	False
2015-01-02	False	True	False	True
2014-12-31	False	False	False	False

	Cases_Nigeria	Cases_Senegal
Date		
2015-01-05	True	True
2015-01-04	True	True
2015-01-03	True	True
2015-01-02	True	True
2014-12-31	True	True

```
[21]: ebola.isnull().sum()
```

Day	0
Cases_Guinea	29
Cases_Liberia	39
Cases_SierraLeone	35
Cases_Nigeria	84
Cases_Senegal	97
Cases_UnitedStates	104
Cases_Spain	106
Cases_Mali	110
Deaths_Guinea	30
Deaths_Liberia	41
Deaths_SierraLeone	35
Deaths_Nigeria	84
Deaths_Senegal	100
Deaths_UnitedStates	104
Deaths_Spain	106
Deaths_Mali	110

dtype: int64

```
[22]: ebola.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 122 entries, 2015-01-05 to 2014-03-22
Data columns (total 17 columns):
Day                122 non-null int64
Cases_Guinea       93 non-null float64
Cases_Liberia      83 non-null float64
Cases_SierraLeone  87 non-null float64
Cases_Nigeria      38 non-null float64
Cases_Senegal      25 non-null float64
Cases_UnitedStates 18 non-null float64
Cases_Spain        16 non-null float64
Cases_Mali         12 non-null float64
Deaths_Guinea      92 non-null float64
Deaths_Liberia     81 non-null float64
Deaths_SierraLeone 87 non-null float64
Deaths_Nigeria     38 non-null float64
```

```
Deaths_Senegal      22 non-null float64
Deaths_UnitedStates 18 non-null float64
Deaths_Spain        16 non-null float64
Deaths_Mali         12 non-null float64
dtypes: float64(16), int64(1)
memory usage: 17.2 KB
```

```
[23]: ebola['Deaths_Mali'].count()
```

```
12
```

Q3-a: Count the number of missing entries in each column of the df DataFrame from Q2.

```
[24]: df.isnull().sum()
```

```
Capacity    2
Demand      1
dtype: int64
```

Q3-b: Write a command to count the number of rows in which the difference between capacity and demand is missing.

```
[25]: (df['Capacity']-df['Demand']).isnull().sum()
```

```
3
```

Q3-c: Similar to Q3-b, except count the number of rows in which the difference is not missing.

```
[26]: (df['Capacity']-df['Demand']).count()
```

```
1
```

```
[27]: (df['Capacity']-df['Demand']).notnull().sum()
```

```
1
```


Filling Missing Values

```
[28]: guinea=ebola['Cases_Guinea']
      guinea.head()
```

```
Date
2015-01-05    2776.0
2015-01-04    2775.0
2015-01-03    2769.0
2015-01-02         NaN
2014-12-31    2730.0
Name: Cases_Guinea, dtype: float64
```

```
[53]: guinea.plot()
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f3124f83208>
```

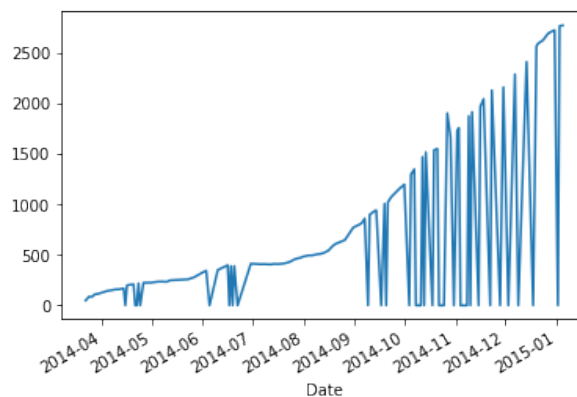


```
[30]: guinea.fillna(0).head()
```

```
Date
2015-01-05    2776.0
2015-01-04    2775.0
2015-01-03    2769.0
2015-01-02         0.0
2014-12-31    2730.0
Name: Cases_Guinea, dtype: float64
```

```
[31]: guinea.fillna(0).plot()
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f312501ee80>
```




```
[32]: guinea.fillna(method='ffill').head()
```

```
Date
2015-01-05    2776.0
2015-01-04    2775.0
2015-01-03    2769.0
2015-01-02    2769.0
2014-12-31    2730.0
Name: Cases_Guinea, dtype: float64
```

```
[33]: guinea.fillna(method='ffill').plot()
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f3124f929b0>
```



```
[34]: guinea.fillna(method='bfill').head()
```

```
Date
2015-01-05    2776.0
2015-01-04    2775.0
2015-01-03    2769.0
2015-01-02    2730.0
2014-12-31    2730.0
Name: Cases_Guinea, dtype: float64
```

```
[35]: guinea.interpolate().head()
```

```
Date
2015-01-05    2776.0
2015-01-04    2775.0
2015-01-03    2769.0
2015-01-02    2749.5
2014-12-31    2730.0
Name: Cases_Guinea, dtype: float64
```

```
[36]: df
```

```
   Capacity  Demand
0      NaN    10.0
1    50.0     NaN
2      NaN    20.0
3    20.0    10.0
```

```
[37]: df['Capacity'].fillna(df['Demand'])
```

```
0    10.0
1    50.0
2    20.0
3    20.0
Name: Capacity, dtype: float64
```

```
[38]: df.dropna()
```

```
   Capacity  Demand
3    20.0    10.0
```

```
[39]: df['Capacity'].first_valid_index()
```

1

```
[40]: df['Capacity'].last_valid_index()
```

3

Q4-a: An analyst would like to calculate the average value of the “Cases_Guinea” column of the ebola Dataset. The analyst runs the below command. Explain why this result might be misleading.

```
[41]: guinea.mean()
```

911.0645161290323

Q4-b: Write a command that corrects the above issue.

```
[42]: guinea.fillna(method='ffill').mean()
```

1023.7295081967213

Any of the following also obtains a similar result.

```
[43]: guinea.fillna(method='bfill').mean()
```

1005.139344262295

```
[44]: guinea.interpolate().mean()
```

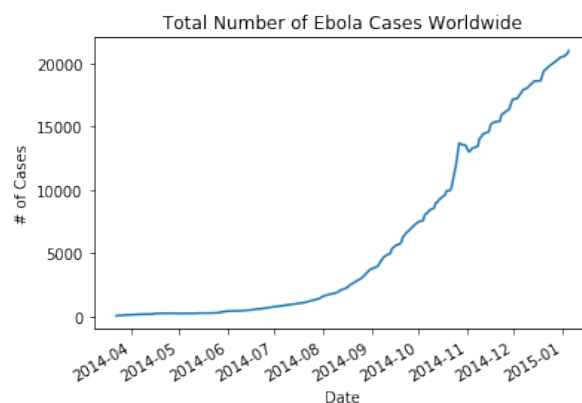
1014.4344262295082

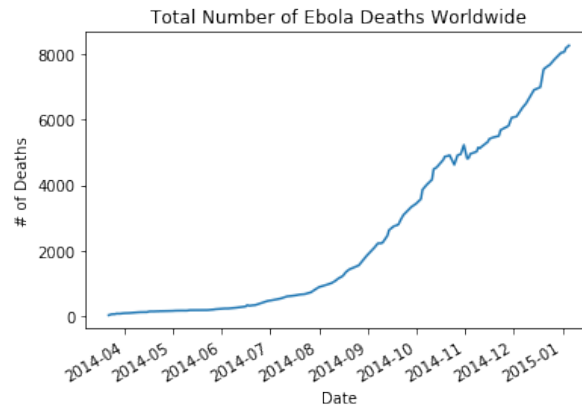
Q5: Plot the total number of cases and the total number of deaths due to Ebola in the data set from all of the countries, while handling missing data appropriately.

```
[45]: cases=ebola.fillna(method='bfill').iloc[:,1:9].sum(axis=1)
      deaths=ebola.fillna(method='bfill').iloc[:,9:].sum(axis=1)
```

```
import matplotlib.pyplot as plt
cases.plot(title='Total Number of Ebola Cases Worldwide')
plt.ylabel('# of Cases')
plt.show()

deaths.plot(title='Total Number of Ebola Deaths Worldwide')
plt.ylabel('# of Deaths')
plt.show()
```





Q6: Obtain the first non-null value of the column “Cases_UnitedStates”, as well as the date on which this is recorded. (Since the dates are ordered backward, this is the last recorded value and the last recorded date). Do the same also for the last non-null value of the column.

```
[46]: # Index of first non-null value
      usa=ebola['Cases_UnitedStates']
      usa.first_valid_index()
```

```
Timestamp('2014-12-07 00:00:00')
```

```
[47]: # First non-null value
      usa[usa.first_valid_index()]
```

```
4.0
```

```
[48]: # Index of last non-null value
      usa.last_valid_index()
```

```
Timestamp('2014-10-01 00:00:00')
```

```
[49]: # Last non-null value
      usa[usa.last_valid_index()]
```

```
1.0
```

Q7: Based on the information in the dataset, obtain an estimate of the number of cases of ebola in each of the countries in the dataset on Jan 5, 2015. Appropriately display the information in a pie chart.


```
[50]: ebola.fillna(method='bfill').iloc[0,1:9]
```

```
Cases_Guinea          2776.0
Cases_Liberia         8166.0
Cases_SierraLeone     10030.0
Cases_Nigeria         20.0
Cases_Senegal          1.0
Cases_UnitedStates     4.0
Cases_Spain            1.0
Cases_Mali             7.0
Name: 2015-01-05 00:00:00, dtype: float64
```

```
[51]: ebola_cleaned=ebola.fillna(method='bfill')
      result=ebola_cleaned.iloc[0,1:4]
      result['Other']=sum(ebola_cleaned.iloc[0,4:9])

      import matplotlib.pyplot as plt
      result.plot(kind='pie',title='Total Ebola Cases by Country as of Jan 5, 2015')
      plt.ylabel('')
      plt.show()
```

Total Ebola Cases by Country as of Jan 5, 2015



This is how you do it for an arbitrary date that may or may not be in the Dataset. The following example is for 2014 Sep 1.

```
[54]: import numpy as np
      ebola.index[np.where(ebola.index<pd.datetime(2014,7,1))].max()
Timestamp('2014-06-30 00:00:00')

[55]: date=ebola.index[np.where(ebola.index<pd.datetime(2014,7,1))].max()
      result=ebola_cleaned.loc[date,:].iloc[1:4]
      result['Other']=sum(ebola_cleaned.loc[date,:].iloc[4:9])

      import matplotlib.pyplot as plt
      result.plot(kind='pie',title='Total Ebola Cases by Country as of July 1, 2015')
      plt.ylabel('')
      plt.show()
```

Total Ebola Cases by Country as of July 1, 2015

