

Session 14: Data Aggregation using Groupby

Setting up

```
[2]: import pandas as pd
      base='https://raw.githubusercontent.com/chendaniely/pandas_for_everyone/master/data/'
      filename='gapminder.tsv'
      gapminder=pd.read_csv(base+filename,sep='\t')
      gapminder.head()
```


	country	continent	year	lifeExp	pop	gdpPercap
0	Afghanistan	Asia	1952	28.801	8425333	779.445314
1	Afghanistan	Asia	1957	30.332	9240934	820.853030
2	Afghanistan	Asia	1962	31.997	10267083	853.100710
3	Afghanistan	Asia	1967	34.020	11537966	836.197138
4	Afghanistan	Asia	1972	36.088	13079460	739.981106

1. Grouping by One Column

```
[2]: gapminder.groupby('year')['lifeExp'].mean().head()
```

```
year
1952    49.057620
1957    51.507401
1962    53.609249
1967    55.678290
1972    57.647386
Name: lifeExp, dtype: float64
```


```
[4]: gapminder.groupby('year')['lifeExp'].mean()\
      .plot(title='Average life expectancy of all countries')
```



```
[4]: gapminder.groupby('year')[['gdpPercap','lifeExp']].mean().head()
```

	gdpPercap	lifeExp
year		
1952	3725.276046	49.057620
1957	4299.408345	51.507401
1962	4725.812342	53.609249
1967	5483.653047	55.678290
1972	6770.082815	57.647386

```
[5]: gapminder.groupby('year')[['gdpPercap', 'lifeExp']].mean()\
      .plot(title='GDP per capita and life expectancy of all countries', subplots=True)
```



List of Pandas methods built into groupby:


Method	Description
count	Number of elements not including NaN values.
size	Number of elements including NaNs.
sum	Total.
mean	Average.
median	Median.
var	Variance.
std	Standard deviation.
min	Minimum.
max	Maximum.
quantile(0.28)	28th percentile from the bottom.
describe	count, mean, std, min, 25%, 50%, 75%, and max
first	first non-null value.
last	last non-null value.
nth(3)	3rd value (does not skip null values).

```
[6]: import numpy as np
      gapminder.groupby('year').agg({'lifeExp': 'mean', 'gdpPercap': 'median'}).head()
```

```
      lifeExp    gdpPercap
year
1952  49.057620  1968.528344
1957  51.507401  2173.220291
1962  53.609249  2335.439533
1967  55.678290  2678.334741
1972  57.647386  3339.129407
```

Q1: Create a bar plot comparing the total populations of each continent in 2007, as below.
(Hint: First use “query” to filter for year being 2007, then group by the continent and compute the sum of the “pop” column. Then sort the result using “sort_values” and plot using “kind=‘bar’”. All this can be chained together into one line.)


```
[7]:
```




Q2: Create the following plots to compare the standard deviation in GDP per capita and life expectancy across countries within each continent, in 1952 and in 2007.

(Hint: for each graph, first use “query” to filter for the year, then group by the continent, and compute the standard deviation of both “gdpPercap” and “lifeExp”. Then plot using “subplots=True”. Each plot can be created using one line by chaining together commands.)

[8] :




[9] :



Q3: Plot the average GDP per capita over the years in the dataset for all countries in the continent “Americas”, as below.

[10]:



2. Grouping by Multiple Columns

Grouping by multiple columns is completely analogous, except for the result having multiple levels of index.

```
[11]: result=gapminder.groupby(['year','continent'])['lifeExp'].mean()
      result.head(10)
```

```
year  continent
1952  Africa      39.135500
      Americas    53.279840
      Asia        46.314394
      Europe      64.408500
      Oceania     69.255000
1957  Africa      41.266346
      Americas    55.960280
      Asia        49.318544
      Europe      66.703067
      Oceania     70.295000
Name: lifeExp, dtype: float64
```

```
[12]: gapminder.groupby(['year','continent']).std().head()
```

```
year  continent  lifeExp  pop  gdpPercap
1952  Africa      5.151581  6.317450e+06  982.952116
      Americas    9.326082  3.234163e+07  3001.727522
      Asia        9.291751  1.132267e+08  18634.890865
      Europe      6.361088  1.724745e+07  3114.060493
      Oceania     0.190919  4.735083e+06  365.560078
```

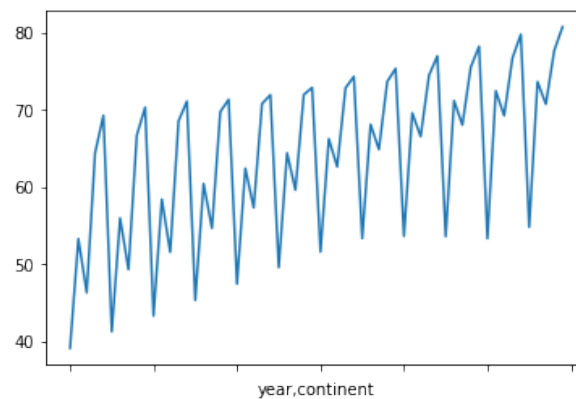
```
[13]: gapminder.groupby(['year', 'continent'])\
      .agg({'lifeExp': 'mean', 'gdpPercap': 'median'}).head()
```

		lifeExp	gdpPercap
year	continent		
1952	Africa	39.135500	987.025569
	Americas	53.279840	3048.302900
	Asia	46.314394	1206.947913
	Europe	64.408500	5142.469716
	Oceania	69.255000	10298.085650

2.1 Working with MultiIndex Objects

The Series “result” has two levels of index, year and continent. This is an example of a Hierarchical Index or MultiIndex in Pandas. The data is the rightmost column above, while the two columns on the left are both row labels.

```
[14]: result.plot()
```



You can change a Series to a DataFrame by the command “reset_index”, which get rid of the MultiIndex.

```
[15]: result.reset_index().head()
```

	year	continent	lifeExp
0	1952	Africa	39.135500
1	1952	Americas	53.279840
2	1952	Asia	46.314394
3	1952	Europe	64.408500
4	1952	Oceania	69.255000

```
[16]: result.reset_index().groupby('continent')['lifeExp'].std()
```

continent	
Africa	5.443393
Americas	6.663477
Asia	8.236002
Europe	4.042574
Oceania	3.843914

Name: lifeExp, dtype: float64

The “unstack” method moves one level of the MultiIndex to the columns, as below. The default behavior is to move the right most level.


```
[17]: result.unstack().head()
```

continent	Africa	Americas	Asia	Europe	Oceania
year					
1952	39.135500	53.27984	46.314394	64.408500	69.255
1957	41.266346	55.96028	49.318544	66.703067	70.295
1962	43.319442	58.39876	51.563223	68.539233	71.085
1967	45.334538	60.41092	54.663640	69.737600	71.310
1972	47.450942	62.39492	57.319269	70.775033	71.910

```
[18]: result.unstack(level=0).iloc[:5,:5]
```

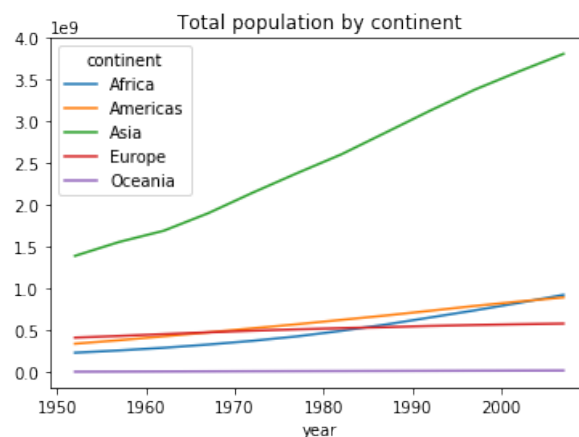
year	1952	1957	1962	1967	1972
continent					
Africa	39.135500	41.266346	43.319442	45.334538	47.450942
Americas	53.279840	55.960280	58.398760	60.410920	62.394920
Asia	46.314394	49.318544	51.563223	54.663640	57.319269
Europe	64.408500	66.703067	68.539233	69.737600	70.775033
Oceania	69.255000	70.295000	71.085000	71.310000	71.910000

```
[19]: gapminder.groupby(['year', 'continent'])['lifeExp'].mean()\
      .unstack().plot(title="Life expectancy by continent")
```



Q4: Plot the trend in total population of each continent as below.


```
[20]:
```



Q5: Compute the average GDP per capita for each continent in 1952 and 2007, and plot the ratio.

(Hint: a quick way is to first group by the continent and year and compute the average GDP Per capita for each combination, then unstack it so that the years are the columns, similar to in Out[30]. Then you can compute the desired ratio by dividing the column for 2007 by the column for 1952.)

[22] :



Q6. Create a plot over time of the difference in total GDP between the richest and the poorest continent, as below.

(Hint: first add a “GDP” column in the gapminder DataFrame by multiplying the “gdpPer-cap” and “pop” columns. Then create a “gdpSum” DataFrame by grouping by the year and continent, and summing the GDPs. See below for what this DataFrame looks like. Using this DataFrame, you can compute a Series called “maxGDP” by grouping by the year and finding the max GDP, and similarly compute a Series called “minGDP”. Both of these are indexed by year. Finally, subtract maxGDP by minGDP and plot the result.)

[23] :

	year	continent	GDP
0	1952	Africa	3.115993e+11
1	1952	Americas	2.943475e+12
2	1952	Asia	1.125160e+12

[26] :

